

Used_Vehicle_Data_Preprocessing

December 5, 2021

Author: Ekin Ugurel

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: # Read raw data
df = pd.read_csv('C:/Users/ekino/Downloads/vehicles.csv')
```

```
[3]: df.columns
```

```
[3]: Index(['id', 'url', 'region', 'region_url', 'price', 'year', 'manufacturer',
          'model', 'condition', 'cylinders', 'fuel', 'odometer', 'title_status',
          'transmission', 'VIN', 'drive', 'size', 'type', 'paint_color',
          'image_url', 'description', 'county', 'state', 'lat', 'long',
          'posting_date'],
          dtype='object')
```

```
[4]: # Statistical Analysis to get a better understanding of database
df.describe()
```

```
[4]:
```

	id	price	year	odometer	county \
count	4.268800e+05	4.268800e+05	425675.000000	4.224800e+05	0.0
mean	7.311487e+09	7.519903e+04	2011.235191	9.804333e+04	NaN
std	4.473170e+06	1.218228e+07	9.452120	2.138815e+05	NaN
min	7.207408e+09	0.000000e+00	1900.000000	0.000000e+00	NaN
25%	7.308143e+09	5.900000e+03	2008.000000	3.770400e+04	NaN
50%	7.312621e+09	1.395000e+04	2013.000000	8.554800e+04	NaN
75%	7.315254e+09	2.648575e+04	2017.000000	1.335425e+05	NaN
max	7.317101e+09	3.736929e+09	2022.000000	1.000000e+07	NaN

	lat	long
count	420331.000000	420331.000000
mean	38.493940	-94.748599
std	5.841533	18.365462
min	-84.122245	-159.827728
25%	34.601900	-111.939847

50%	39.150100	-88.432600
75%	42.398900	-80.832039
max	82.390818	173.885502

```
[5]: # Check for missing values
null_count = pd.DataFrame({'Null': df.isnull().sum()})
# Check for percent of values missing
length=len(df)
percent_null = round((null_count['Null']/length)*100,1)
null_count['Percentage'] = percent_null
# Sort from highest percentage to lowest
null_count.sort_values(by='Null', ascending=False)
```

```
[5]:
```

	Null	Percentage
county	426880	100.0
size	306361	71.8
cylinders	177678	41.6
condition	174104	40.8
VIN	161042	37.7
drive	130567	30.6
paint_color	130203	30.5
type	92858	21.8
manufacturer	17646	4.1
title_status	8242	1.9
lat	6549	1.5
long	6549	1.5
model	5277	1.2
odometer	4400	1.0
fuel	3013	0.7
transmission	2556	0.6
year	1205	0.3
description	70	0.0
image_url	68	0.0
posting_date	68	0.0
url	0	0.0
price	0	0.0
state	0	0.0
region_url	0	0.0
region	0	0.0
id	0	0.0

```
[6]: # Drop columns that have too many missing values or that are irrelevant
df.drop(['posting_date',
        ↪ 'county', 'VIN', 'url', 'region_url', 'image_url', 'id', 'lat', 'long', 'description'],
        ↪ axis=1, inplace=True)
df.shape
```

[6]: (426880, 16)

```
[7]: # Check for duplicates
df.duplicated().sum()
```

[7]: 56415

```
[8]: # Drop duplicates and keep one of each
df = df.drop_duplicates(keep='first')
```

```
[9]: # remove inconsistent data entry (e.g. spaces in the cell)
columns = ['manufacturer', 'condition', 'cylinders', 'fuel', 'title_status', 'transmission', 'drive', 'size', 'type', 'paint_color']

for i in columns:
    df[i] = df[i].str.strip()
```


```
[11]: # Drop 10% of each side on price (outliers)
sort = sorted(df['price'])
q1, q2 = np.percentile(sort, [10,90])
print(q1, q2)
```

1200.0 37740.0

```
[12]: df = df[(df.price <= 37740.0) & (df.price >= 1200)]
df.shape
```

[12]: (296798, 16)


```
[14]: # Visualize odometer distribution to check for outliers
plt.figure(figsize=[7,5])
odo = sns.scatterplot(x = df['odometer'], y=df['price'])
```




```
[15]: # Since cars after 300,000 miles of use become almost obsolete, drop all values
      ↪ greater than 3e5
df = df[(df.odometer < 3e5)]
```

```
[16]: # Also drop odometers that equal 0, as we are concerned with used cars (not new)
df.drop(df[df['odometer']==0.0].index, inplace=True)
```

```
[17]: plt.figure(figsize=[7,5])
odo = sns.scatterplot(x = df['odometer'], y=df['price'])
```




```
[18]: # handle outliers for "year"
plt.figure(figsize=[7,5])
odo = sns.scatterplot(x = df['year'], y=df['price'])
```



```
[19]: # Since outliers begin (roughly) before the year 1945, drop all such entries
df.drop(df[df['year'] <= 1945].index, inplace=True)
```

```
[20]: # Check the scatter plot again
plt.figure(figsize=[7,5])
odo = sns.scatterplot(x = df['year'], y=df['price'])
```



```
[21]: plt.figure(figsize=[10,5])
plt.subplot(121)
sns.distplot(df['price'], bins = 5)
plt.subplot(122)
sns.distplot(df['odometer'], bins=5)
```


C:\Users\ekino\anaconda3\lib\site-packages\seaborn\distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

C:\Users\ekino\anaconda3\lib\site-packages\seaborn\distributions.py:2619:
FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

```
[21]: <AxesSubplot:xlabel='odometer', ylabel='Density'>
```

```
[22]: # Fill in NaN values for columns "condition" and "title status"
bins = [0,30000,60000,90000,115000,150000,1000000]
groups = df.groupby(['title_status', pd.cut(df.odometer,bins)])
groups.size().unstack()
```

```
[22]: odometer      (0, 30000]  (30000, 60000]  (60000, 90000]  (90000, 115000]  \
title_status
clean           47388           42932           42568           39304
lien             127             166             261             192
missing          130              53              63              85
parts only         14              12               7              11
rebuilt           981            1393            1424             898
salvage           429             535             594             446

odometer      (115000, 150000]  (150000, 1000000]
title_status
clean           47848           55974
lien             210             204
missing           58              78
parts only         11              23
rebuilt            896             716
salvage           542             730
```

```
[23]: # Since an overwhelming majority of the entries have "clean" titles, this
      ↪ column becomes insignificant. So we drop this as well.
df.drop(['title_status'], axis = 1, inplace=True)
```



```
[24]: # Now, let's check the distribution for condition
bins = [0,30000,60000,90000,115000,150000,1000000]
groups = df.groupby(['condition', pd.cut(df.odometer,bins)])
groups.size().unstack()
```

```
[24]: odometer    (0, 30000]  (30000, 60000]  (60000, 90000]  (90000, 115000]  \
condition
excellent          4606             8272             13169             14048
fair                252              188              348              583
good              30487            20022            12201            8393
like new           2719             2670             2692            2374
new                233              69              82              53
salvage            29              40              38              50

odometer    (115000, 150000]  (150000, 1000000]
condition
excellent          16917             15101
fair               948              2916
good             10928            17649
like new          2563             1839
new               75              74
salvage           63              135
```

```
[25]: # Replace missing entries with the median condition for each odometer level
g1 = (df['odometer'] > 60000) & (df['odometer'] <= 150000)
g2 = (df['odometer'] <= 60000) | (df['odometer'] > 150000)

df.loc[g1,'condition']=df.loc[g1,'condition'].fillna('excellent')
df.loc[g2,'condition']=df.loc[g2,'condition'].fillna('good')
```

```
[26]: # Check for missing values again
null_count = pd.DataFrame({'Null': df.isnull().sum()})
# Check for percent of values missing
length=len(df)
percent_null = round((null_count['Null']/length)*100,1)
null_count['Percentage'] = percent_null
# Sort from highest percentage to lowest
null_count.sort_values(by='Null', ascending=False)
```

```
[26]:
```

	Null	Percentage
size	203732	69.8
cylinders	115143	39.5
drive	88694	30.4
paint_color	80567	27.6
type	62264	21.3
manufacturer	10519	3.6
model	3076	1.1

fuel	1578	0.5
transmission	1078	0.4
year	471	0.2
region	0	0.0
price	0	0.0
condition	0	0.0
odometer	0	0.0
state	0	0.0

```
[27]: # Drop 'size' column, as it has too many missing values
df.drop('size', axis=1, inplace=True)
# For columns with less than 4% missing values, drop the empty rows
df = df.dropna(subset=['manufacturer', 'model', 'fuel', 'transmission', 'year'])
```

```
[28]: # Use "ffill" to propagate non-null values forward or backward for columns with
      ↪ 10-30% missing values
columns = ['drive', 'type', 'paint_color']
for i in columns:
    df[i] = df[i].fillna(method='ffill')
```

C:\Users\ekino\AppData\Local\Temp\ipykernel_52636\4242946631.py:4:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df[i] = df[i].fillna(method='ffill')
```

```
[29]: # Drop the empty rows if there are any left
df = df.dropna(subset=['type'])
df = df.dropna(subset=['drive'])
```

```
[30]: # Check for missing values again
null_count = pd.DataFrame({'Null': df.isnull().sum()})
# Check for percent of values missing
total=len(df)
percent_null = round((null_count['Null']/total)*100,1)
null_count['Percentage'] = percent_null
# Sort from highest percentage to lowest
null_count.sort_values(by='Null', ascending=False)
```

```
[30]:
```

	Null	Percentage
cylinders	107802	39.1
region	0	0.0
price	0	0.0
year	0	0.0
manufacturer	0	0.0

model	0	0.0
condition	0	0.0
fuel	0	0.0
odometer	0	0.0
transmission	0	0.0
drive	0	0.0
type	0	0.0
paint_color	0	0.0
state	0	0.0

```
[31]: # Only column left to deal with is 'cylinders'.
# We can use the 'drive' column to make best-guesses about the missing values
      ↪ in 'cylinder'
df.groupby(['drive', 'cylinders']).cylinders.count()
```

```
[31]: drive  cylinders
4wd    10 cylinders    371
      12 cylinders      8
      3 cylinders    108
      4 cylinders  15361
      5 cylinders    335
      6 cylinders  30062
      8 cylinders  23839
      other         177
fwd    10 cylinders     38
      12 cylinders      4
      3 cylinders    247
      4 cylinders  35256
      5 cylinders    791
      6 cylinders  21727
      8 cylinders    3031
      other         245
rwd    10 cylinders    464
      12 cylinders     52
      3 cylinders     25
      4 cylinders   5031
      5 cylinders    134
      6 cylinders  13983
      8 cylinders  16506
      other         157
Name: cylinders, dtype: int64
```

```
[32]: # Fill in the median value of "cylinders" for each type of "drive"
values = {'4wd': '6 cylinders', 'fwd': '4 cylinders', 'rwd': '8 cylinders'}
df.loc[df['cylinders'].isna(), 'cylinders'] = df.loc[df['cylinders'].
      ↪ isna(), 'drive'].map(lambda x: values[x])
```

```
[33]: # Check for missing values again
null_count = pd.DataFrame({'Null': df.isnull().sum()})
# Check for percent of values missing
total=len(df)
percent_null = round((null_count['Null']/total)*100,1)
null_count['Percentage'] = percent_null
# Sort from highest percentage to lowest
null_count.sort_values(by='Null', ascending=False)
```

```
[33]:
```

	Null	Percentage
region	0	0.0
price	0	0.0
year	0	0.0
manufacturer	0	0.0
model	0	0.0
condition	0	0.0
cylinders	0	0.0
fuel	0	0.0
odometer	0	0.0
transmission	0	0.0
drive	0	0.0
type	0	0.0
paint_color	0	0.0
state	0	0.0

```
[34]: # Since 'region' and 'state' are directly related to one another, join these
      ↪ two columns
df['region'] = df['region'] + ' (' + df['state'] + ')'
df.drop(['state'], axis = 1, inplace=True)
```

```
[35]: df.shape
```

```
[35]: (275754, 13)
```

```
[40]: # change the order of columns such that 'price' is last
columns = ['region', 'year', 'manufacturer', 'model', 'condition',
           'cylinders', 'fuel', 'odometer',
           ↪ 'transmission', 'drive', 'type', 'paint_color', 'price']
df = df.reindex(columns=columns)
df.head(1)
```

```
[40]:
```

	region	year	manufacturer	model	condition	cylinders	fuel	\
31	auburn (al)	2013.0	ford	f-150 xlt	excellent	6 cylinders	gas	
	odometer	transmission	drive	type	paint_color	price		
31	128000.0	automatic	rwd	truck	black	15000		

```
[37]: df.describe()
```

```
[37]:
```

	year	odometer	price
count	275754.000000	275754.000000	275754.000000
mean	2010.922514	96569.286777	16157.680465
std	8.243936	61415.561154	10013.385017
min	1946.000000	1.000000	1200.000000
25%	2008.000000	43213.000000	7450.000000
50%	2013.000000	93000.000000	14500.000000
75%	2016.000000	139594.000000	23999.000000
max	2022.000000	299999.000000	37740.000000

```
[41]: # This dataset looks good. We will download this as a CSV and use it in our
      ↪analysis.
df.to_csv('preprocessedusedcars.csv')
```

```
[42]: # However, we also want to check if there are highly correlated features using
      ↪a Correlation Matrix.
# If so, we would drop one of these features so that our analysis isn't
      ↪redundant
df2 = df.copy() # make a copy of dataframe
```

```
[44]: # In order to check for correlation, we must convert all "categorical" features
      ↪into numerical features
# We use the LabelEncoder function from sklearn to do this
from sklearn.preprocessing import LabelEncoder
categories = ['region', 'manufacturer', 'model', 'cylinders', 'fuel',
      ↪'transmission',
      ↪'drive', 'type', 'paint_color', 'condition']
encoder = LabelEncoder()
encoded = df[categories].apply(encoder.fit_transform)
df2.drop(categories, axis=1, inplace=True)
df2 = pd.concat([encoded, df2], axis=1)

df2.head(10)
```


```
[44]:
```

	region	manufacturer	model	cylinders	fuel	transmission	drive	type	\
31	18	13	7661	5	2	0	2	10	
32	18	14	15026	6	2	2	0	8	
33	18	7	15203	5	2	2	0	8	
34	18	38	16252	5	2	0	0	10	
35	18	7	4875	5	2	2	0	8	
37	18	20	4381	5	2	0	0	8	
38	18	20	17954	5	2	2	0	7	
39	18	7	15257	5	2	2	0	8	
40	18	7	4871	5	4	2	0	8	
41	18	38	16289	5	4	2	0	8	

	paint_color	condition	year	odometer	price
31	0	0	2013.0	128000.0	15000
32	0	2	2012.0	68696.0	27990
33	9	2	2016.0	29499.0	34590
34	5	0	2019.0	43000.0	35000
35	8	2	2016.0	17302.0	29990
37	8	0	1992.0	192000.0	4500
38	9	2	2017.0	30041.0	32990
39	10	2	2017.0	40784.0	24590
40	1	2	2016.0	34940.0	30990
41	8	2	2014.0	17805.0	27990

```
[52]: # Correlation Heat Map
sns.set(style='whitegrid')
cor = df2.drop(columns=['price']).corr()
mask = np.zeros_like(cor, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True
f, ax = plt.subplots(figsize=(10,10))
cmap = sns.diverging_palette(220, 10, as_cmap=True)
sns.heatmap(cor, mask=mask, cmap=cmap, vmax=.3, center=0, square=True,
↪linewidths=.5, cbar_kws={"shrink":.5}, annot=True)
```

```
[52]: <AxesSubplot:>
```



```
[53]: # No features are highly correlated, so we keep the set as is.
df2.to_csv('numericalpreprocesseddata.csv')
```